

→ Bagging

- What is Bagging?: An ensemble learning technique that trains multiple models independently on different bootstrap samples and combines their predictions to improve stability and accuracy.
- Purpose (When): To reduce variance and overfitting in high variance models (especially Decision Trees).
- Why it is used
 - Stabilizes unstable models
 - Reduces overfitting
 - Improves generalization
 - Parallelizable
- How it works (Algorithm)
 - Create multiple bootstrap samples (sampling with replacement)
 - Train a base learner on each sample
 - Make predictions independently
 - Aggregate results: Classification → majority vote
Regression → mean
- Formula
 - Classification: $\hat{y} = \text{mode}(h_1(x), h_2(x), \dots, h_n(x))$
 - Regression: $\hat{y} = \frac{1}{n} \sum h_i(x)$

- Technical Details

- Ensemble type: Parallel
- Reduces variance, not bias
- Base models are independent of each other

- Parameters: Number of estimators, base estimator type (gini/entropy), bootstrap size (depth)

- ## - Pros
- Reduces overfitting
 - Simple and effective
 - Parallel execution

- Cons

- No bias reduction
- Less interpretable
- Higher computation time

- Real World Applications

- Financial risk models
- Medical prediction systems
- Baseline ensemble models.

→ Gradient Boosting

- What is Gradient Boosting

Gradient Boosting is a sequential ensemble technique where each new model corrects the errors of the previous models using gradient descent.

- Purpose ~~(why)~~ (When)

To reduce bias and build strong predictive models from weak learners.

- Why it is used

- Handles complex non-linear patterns
- High predictive accuracy
- Flexible loss function

- How it works (Algo)

- Start with a simple model (initial prediction)
- Compute residual errors
- Train next model on residuals
- Add model to ensemble with learning rate
- Repeat sequentially

- Formula

$$F_m(x) = F_{m-1}(x) + n \cdot h_m(x) \quad ; \quad n = \text{learning rate}$$

$\therefore h_m = \text{weak learner}$

- Technical Details

- Ensemble type: Sequential
- Optimizes arbitrary loss functions
- Uses gradient descent in function space

- Parameters

- Learning rate
- Number of estimators
- Tree depth
- Loss function

- Pros

- High accuracy
- Handles complex data
- Custom loss functions

- Cons

- Slow training with many estimators
- Sensitive to hyperparameters
- Prone to overfitting if not tuned

- Real World Applications (where)

- Fraud detection
- Ranking systems
- Credit scoring

→ AdaBoost (Adaptive Boosting)

- What is AdaBoost

~~Ada~~ AdaBoost is a boosting algorithm that focuses on misclassified samples by increasing their weight in ~~seqt~~ subsequent models.

- Purpose (When)

To convert weak learners into a strong classifier by focusing on hard examples.

- Why it is used

- Simple boosting algorithm
- Strong theoretical foundation
- Works well with weak learners

- How it works (Algo)

- Assign equal weights to all samples
- Train weak learner
- Increase weights of misclassified points
- Train next learner on reweighted data
- Combine learners using weighted voting.

- Formula

$$\text{Final model: } F(x) = \sum \alpha_m h_m(x)$$

$$\text{Learner weight: } \alpha_m = \frac{1}{2} \ln \frac{1 - \text{error}}{\text{error}}$$

- Technical Details

- Sequential learning
- Sensitive to noisy data
- Uses exponential loss

- Parameters

- Number of estimators
- Learning rate
- Base estimator

- Pros

- Improves weak models
- Simple implementation
- Less overfitting than trees

- Cons

- Sensitive to noise & outliers
- Poor performance on complex data
- Not scalable

- Real World Applications (Where)

- Face detection
- Text classifications
- Early ML systems

XGBoost (Extreme Gradient Boost)

- What is XGBoost

XGBoost is an optimized, scalable implementation of Gradient Boosting designed for speed, performance and regularization.

- Purpose (When)

To achieve state of the art accuracy on structured/tabular data.

- Why it is used

- Extremely fast
- Built in regularization
- Handles missing values
- Winner of many ML competitions

- How it works (Alg.)

- Builds trees sequentially
- Uses second-order derivatives (Hessian)
- Applies regularization to trees
- Prunes trees automatically
- Parallelizes tree construction

- Formula

Objective Function: $\text{Obj} = \sum l(y_i, \hat{y}_i) + \sum \Omega(f_k)$

Regularization: $\Omega(f) = \gamma T + \frac{1}{2} \lambda \sum w^2$

- Technical Details

- Uses ~~newton~~ boosting

- Supports sparse data

- Cache-aware and parallel

- Parameters

- n_estimators : no of trees

- learning_rate : step size

- max_depth : tree depth

- subsample : row sampling

- colsample_bytree : Feature sampling

- lambda, alpha : regularization

- Pros

- Very high accuracy

- Built in regularization

- Handles missing values

- Scalable

- Cons

- Complex tuning

- Less interpretable

- Overkill for small datasets

- Real World Applications (Where)

- Kaggle competitions

- Finance risk models

- Search ranking

- Recommendation systems